

Pryzm — Expert Feedback Strategic Playbook

Source: Conversation with senior data/AI expert (Amex, Amazon, Amazon Flex Luxembourg, Colgate-Palmolive Times Award winner, L'Oréal AI Commerce Studio lead, advised by Jeremy — 8+ startup exec, 90K+ followers, worked with Bosch / Oatly / US & Dubai startups) **Subject:** Diagnosis of current Pryzm / Scherzinger dashboard + playbook to make it corporate-ready **Status:** Strategic — must inform Phase 2+ design before next client-facing iteration

0. TL;DR — The One-Line Diagnosis

▮ "Value-wise there is so much. Structure-wise there is nothing."

You have built a **fancy Excel with capabilities**. Corporates don't buy capabilities. They buy **financial value, strategic value, and a story their manager can forward up the chain**. Everything below is how to make that pivot — without rebuilding from scratch, because the underlying value is already there. It just isn't **shown**.

1. Who Is Giving This Feedback (Credibility Context)

Context for why you should weight this feedback heavily — the expert has shipped the exact kind of enterprise AI products Pryzm is trying to be:

- **American Express** — cost optimization data model; tracked compliance rates week-on-week on hotel/flight booking patterns
- **EdTech company** — built an e-learning platform
- **Amazon** — ML product selection algorithm; picked 400 attributes out of 4,000 per product to identify high-potential SKUs for new marketplace launches. Scaled

pipeline from **1M → 50M products/week** for localization

- **Amazon Flex (Luxembourg)** — NLP sentiment pipeline tracking **600K delivery partners across 18 countries**. Extracted satisfaction index across lifecycle stages (onboarding, on-road support, navigation, customer handoff, payout, offboarding). Identified which of **560 microservices** in the Flex app were underperforming
- **Amazon Flex (Part 2)** — built anomaly detection model that reduced bug detection + fix time from **4.5 days → 3 hours** by benchmarking last-30-day feedback averages per problem type, firing alarm on >25% deviation, and auto-routing to the right team via a responsibility matrix
- **Colgate-Palmolive** — **65% ROAS increase, Times Award UK winner**. Vector database matching user attributes to persona attributes most likely to buy each product. Covered audience targeting, bid optimization, budget allocation, and creative optimization (e.g., discovered toothpaste buyers drink a lot of coffee → put coffee cups in ad creative → subconscious connection drives attention)
- **L'Oréal (current project)** — AI Commerce Studio auto-generating content for 60,000 products across 30 e-retailers. SEO + GEO (Generative Engine Optimization for Rufus-style AI shopping assistants). Covers ad creative management, campaign intelligence, bid/visibility monitoring, performance analytics, workflow automation
- **Unilever/Dove budget allocator (latest)** — solves "Dove gives me annual targets across 30 e-retailers in 20 countries with a fixed budget — what can I achieve?" across **16 advertising media** (Prime Video, Twitch, Prime Music, NY Times, sponsored brands, sponsored products, etc.). Uses linear model to find curve flattening + opportunity size per medium, ML-learned weights for optimal allocation, and added a **vector similarity filter** for new products — after Samsung's new phone failed despite matching all prior attributes, they filter training data based on vector resemblance to historical products before predicting

Relevance to Pryzm: Every single one of these projects was about turning raw signal → action → tracked business outcome. That is exactly what Pryzm is missing.

2. What Pryzm Currently Is (Honest Assessment)

2.1 What's built and working

- **Data foundation (strong):** 5,565 invoice records + 4,605 quote records, 89.3% Auftrag-field linkage, data quality flags, Deckungsbeitrag analysis, win-rate by deal size, margin-gap-by-year
- **Forecasting module (very strong — this is the crown jewel):** Monthly drill-down, conservative/primary/aggressive bands, <5% error validated on Q1 2025, walk-forward methodology, Monte Carlo simulations, seasonality detection with web-context enrichment, Pareto-based repeat customer modeling
- **Model-routing agent (strong IP):** Agent trains every model every day on new data, picks the best one per period, produces prediction bands — "this is the biggest mode in the prediction"
- **Margin intelligence:** DB2 erosion 63.6% → 60.6%, 3.9pp quoted-vs-actual gap widening to 4.6pp, BKAGG underperformance, article 200832-E negative margin detection
- **Pricing & Quotes page:** SKU-level pipeline, reactive + prediction data, price sensitivity — combining these at one point gives "current price → what should it be next year"
- **AI Insights page:** "Build a retention plan for this customer" — recomputed nightly with full data access
- **Guardrail / Live Margin Alerts feature:** When a measure drops below threshold, live alert appears for people with access; sales-team deviation triggers "why did you do

this?" accountability flow

- **Forecasting-model feedback:** Scherzinger already confirmed forecast accuracy in meeting; Avana (Indian medical distributor, ~30% YoY growth) validated forecast match with <2% error on predicted-vs-actual 2025 prices — some SKUs matched exactly, some had already been revised even higher than forecast, giving directional confirmation

2.2 What it is, brutally

"What you're doing is, basically, whatever sense the AI can make out of data, you are just throwing it on their face."

- **10 screens × ~30 KPIs = ~300 data points** thrown at whoever logs in
- **No defined persona → no owner → no champion inside the client**
- **No niche** — Pryzm tells the client "here's everything, you figure out what to do"
- **The strongest feature (forecasting) is buried mid-list** — "90% of people will never open it"
- **The unique proposition sits at the bottom of a table** — the one thing only Pryzm can deliver is the least visible thing in the product
- **AI Summary reads like "a college intern's review"** — states what the AI sees, not what a 10-year-experienced pricing manager would take away
- **Missing the corporate loop entirely** — no way for a user to show their manager what they accomplished this week

3. The Core Diagnosis — Raw Signal vs Metrics vs Value

The expert's mental model of what clients see in a data product:

Layer	Example	Audience reaction
Raw signal	"Here are 5,565 invoice records"	Noise. Don't care.
Metrics	"DB2 margin is 60.6%, down 3pp YoY"	Still noise. They already know this — it's in their Excel.
Financial value	"Our predictions saved you €X this quarter / will unlock €600K additional revenue this year"	This is what they want.
Strategic value	"In the next 12 months, our system ensures you never have inventory overload on more than 10% of products"	This is why they sign the renewal.

What Pryzm shows today

- **Good news section:** raw signal ✗
- **KPIs everywhere:** metrics ✗
- **Financial value headline:** missing
- **Strategic value proposition:** missing from first view

Direct quote capturing the principle

"What they want to see instead is financial value and strategic value. Financial value is, how much money did this accuracy increase save for them this year? Or how much is it going to save this year? What did it avoid? What, your prediction, was able to help them already be prepared for. And all of them, converted into their matrix. And therefore, your first page should only be the matrix that are key matrix for this client."

"It should always work backwards from each client. Ask them, what are the 10 things that you always keep in mind? And create them first."

The "unique data points only you give" rule

"After that, the second thing, it should be the unique data points that only you give. Rest all are sheets."

Implication for Pryzm: The forecasting band with <5% error, the SKU-level margin leakage detection, the cross-customer same-volume price discrepancy — these are the unique data points. They belong on page 1, not buried.

4. The Seven Specific Drawbacks (Ranked by Severity)

Drawback 1 — No persona targeting (SEVERITY: HIGHEST)

Problem: One company told the expert they were being "transparent" and gave dashboard access to all teams. He called it "the worst mistake" — and says Pryzm is making the same one.

"Why don't they know, who to give this to? Because you don't know, what you want to show. That is true. You have thrown everything, in their mouth. We don't know, for whom this is important... No one will open. No one will use it. Your product will fail."

Pricing Manager doesn't care about: order recency, product count, quote win rate (that's for sales team)

A Pricing Manager DOES care about: pricing and controlling together, margin leakage per SKU, guardrails, competitor benchmarks

Required fix: Persona-based login. MD sees one view. CFO sees another. Pricing Manager another. Controller another. Sales another. Inventory another. Strategy another.

Drawback 2 — No action orientation (SEVERITY: HIGHEST, this is the game-changer)

Problem: Corporate logic the product ignores.

"How does corporate work? I have to go to my manager and tell him what I have done this week. That manager will go to his manager and tell him what my team has done this week. Now, what you are doing — it is just that, I am giving you all the data, see in this and think what you have to do."

Required fix: See Section 5 — **the Action Dashboard + Feedback Loop**. This is the single biggest idea in the entire conversation.

Drawback 3 — Unique proposition is hidden (SEVERITY: HIGH)

Problem: Forecasting is mid-list. The lowest table on the forecasting page is the unique proposition.

"Your KPIs are very random. Your forecasting thing, it is in the middle of the list, which 90% of people will never open. And even in that, the lowest table is your unique proposition."

Required fix: Whatever is ONLY Pryzm can do (forecasting with agent-routed models, margin leakage cross-customer, SKU-level price recommendation bands) must be on first view, labelled as the proprietary delta.

Drawback 4 — AI Summary quality (SEVERITY: HIGH)

Problem:

"Your AI summary is now the review of a college intern. What he is doing is, he says whatever he sees. AI summary should say: a 10-year-experience pricing manager, what will he take away from this?"

Required fix: Rewrite AI Summary prompt to simulate a domain-expert take-away per persona. Output = synthesis + recommendation + so-what, not description.

Drawback 5 — Over-reliance on buyer (SEVERITY: MEDIUM-HIGH)

Problem:

■ **"Apart from the owner, no user works that hard. And the owner is never a user."**

The person buying Pryzm is NOT the person using it daily. If the daily user has to drill through 10 screens to find their action, they won't. Product dies after the demo.

Required fix: Daily user must open Pryzm and see, within 5 seconds: *"Here are my three actions this week."*

Drawback 6 — Sales team needs simplification (SEVERITY: MEDIUM)

Problem: Sales team at a large Mittelstand manufacturer only needs a traffic-light signal on each quote: green / yellow / red against guardrails set by pricing.

■ **"In such big companies, hands grow, one point they have only one, green, red, and yellow signal."**

Required fix: Separate sales view — no dashboard, just quote-level RAG (Red/Amber/Green) vs pricing guardrails.

Drawback 7 — No internal politics ammunition (SEVERITY: MEDIUM)

Problem: The manager using Pryzm has no weekly artifact to forward upward. That means zero internal champion, zero renewal motive.

Required fix: Auto-generated weekly email: *"Following Pryzm's recommendations, your team captured €X this week. €Y is still open from 24 un-actioned items."* Manager forwards to his manager. His manager asks questions. Pryzm becomes the source of truth.

5. THE BIG IDEA — Action Dashboard + Feedback Loop

This is the most important section in this entire document. Read it twice.

5.1 First View = Action Dashboard (not KPI dashboard)

Replace the current landing page with:

ACTION DASHBOARD – [Persona: Pricing Manager]	
GAINED this month from accepted actions: € 147,300	
LOST this month from un-actioned items: € 82,400	
OPEN opportunity still available: € 296,000	
YOUR ACTIONS THIS WEEK (ranked by impact)	
1. Article 205415-B → raise price from €4.10 → €4.38	
Impact: €12,400 / quarter Customer: 101580 [Assign ▼]	
[✓ Accept & Implement] [✓ Accept, Not Implemented]	
[▲ Accept, Partial: ____€] [× Reject + reason]	
2. Customer 101580 churn risk – retention offer suggested	
[Same 4 buttons]	
... (8 more, ranked)	
CUMULATIVE (last 90 days)	
Recommendations made:	142
Accepted & Implemented:	47 → €421K captured
Accepted, not implemented:	28 → €187K missed
Partially implemented:	19 → €68K of €124K
Rejected:	48

5.2 Four-option feedback loop per recommendation

Every AI-generated recommendation gets these four buttons:

1. **Accepted & Implemented** — user executed as suggested
2. **Accepted, Not Implemented** — user agreed but couldn't execute (resource / timing / blocker — collect reason)
3. **Accepted, Partially Implemented** — user executed but modified (e.g., Pryzm said raise to €780, they raised to €700 — capture the actual value)
4. **Rejected** — user disagreed (collect reason: customer context, market knowledge, strategic reason)

5.3 Why this is transformational

Before: Client asks "did Pryzm help?" → nobody knows. Revenue moved. Was it us? Was it the market? → renewal in doubt.

After: Every week Pryzm generates a clean attribution:

- **"Following our suggestions, you achieved €X cumulatively."** → manager forwards this to his manager ↑
- **"Suggestions accepted but not implemented cost you €Y."** → manager's manager asks "why didn't you implement?" → pressure flows DOWN to use Pryzm more
- **Corporate feedback loop is now running for you, not against you.**

5.4 Assignment / delegation

- **With CRM access (Zoho/SAP):** Pryzm auto-routes actions to the right person based on role/team

- **Without CRM access (minimum viable):** dropdown picker on each action → manager assigns → assignee sees their queue → accepts / declines / comments

If assignee marks "not going to do it" → manager sees the projected €-loss of not acting → internal escalation handles itself.

5.5 Expert's direct quote on why this matters

"That's how corporate works. I gave you 30 actions. You did only 6. From that 6, you got this. But do you know what you have lost from 24?

This is what I wanted. Because this thing — already most of them told — they have in Excel. They have built a lot of things like this in Excel. But the main thing which they don't have is, because we are training the models based on their data. GPU is trained, everything is trained. That they have, they can't implement on Excel because the formula runs."

The moat isn't the metrics — it's the action + accountability + attribution loop, sitting on top of models they can't run in Excel.

6. Persona-Based Views (Design Spec)

Required personas (minimum viable set)

Persona	First-view emphasis	Secondary
MD / CEO	Strategic value proposition headline ("next 12 mo: inventory overload <10%"); financial value captured YTD; 3 red-flag alerts	Drill-down paths, not detail

Persona	First-view emphasis	Secondary
CFO / Controller	Margin trajectory, DB2 gap quoted-vs-actual, forecast band for quarter-end revenue, working capital tied in under-priced SKUs	Customer-level P&L, cost variance
Pricing Manager	Action dashboard (section 5); SKU margin leakage list; cross-customer price-discrepancy table; guardrail breaches	Elasticity where available; competitor PA/PR rejection patterns
Sales Lead	Quote pipeline funnel, win-rate by segment, lost-opportunity value, guardrail breaches by rep	Customer churn risk
Sales Rep (individual)	Traffic-light on active quotes (green/amber/red vs guardrail), next-best-action per open deal	—
Inventory / Supply	Over/under-stock alerts, slow-mover flags, warehouse-held unrealized revenue (the Indian medical distributor found this the most valuable insight — inventory they didn't even know they had)	—

Hierarchy behavior (the MD-emails-down pattern)

Expected user flow:

1. MD opens dashboard → sees alert / red metric
2. MD emails Pricing Manager: *"Are you doing something about this?"*
3. Pricing Manager opens their view → must be able to drill down to that specific issue
4. Pricing Manager accepts/rejects/implements → feedback loop closes

The AI Summary — rewrite

Today it says *what it sees*. Required behavior: simulate a 10-year senior pricing manager writing a Monday-morning summary to leadership.

Prompt pattern (roughly):

"You are a pricing manager with 10 years of experience at a precision manufacturer. Given this week's data, write the 3-paragraph take-away your CFO would want. Lead with financial impact. Flag the one thing that needs attention. End with 'so-what' — what changes next week."

7. Traffic Light for Sales (Simplification Layer)

Mittelstand sales teams don't want a dashboard. They want a signal.

Proposed sales interface:

Quote #12847 – Customer: 102330 – Article: 201104-G



GREEN

Margin 34% – above guardrail

Proceed with confidence

Quote #12848 – Customer: 101580 – Article: 200832-E



RED

Margin 8% – below guardrail

Requires pricing-team approval

[Escalate]

- Guardrails set by Pricing Manager
- Sales works within them
- Any deviation triggers automatic visibility to team leader → "why did you do this?" accountability
- Integrates with existing Live Margin Alerts feature (which is already built — just needs to be the sales team's primary interface, not a buried alert)

8. Transparency & Risk Mitigation Playbook

This is a professional-maturity layer every serious enterprise AI product must have. Pryzm doesn't yet.

8.1 The standard fine-print (already covered)

"These tools are written in their fine print that you are using AI and these are only recommendations. Please use your judgment and business knowledge to take these calls."

This protects you legally but **does not build trust**. Trust requires the next layer.

8.2 The trust-building layer (missing)

When a Pryzm recommendation backfires, Pryzm must be the FIRST to flag it.

Expert's exact script:

"We just realized this month — our recommendation of increasing inventory / increasing the price — backfired. We immediately noticed it in the last two weeks and therefore we are suggesting the rollback of this increase in price. Or please decide what call you want to take. This is an alarming situation.
Along with this, you have to say — we also looked deeper into why this would have happened. And these are the steps that we are taking for this."

8.3 Why this works

- Business owners know 50% of AI suggestions will be wrong. They expect it.
- The damage isn't the wrong call — it's **discovering the wrong call late**
- If Pryzm is the fastest source of truth for both good AND bad outcomes, Pryzm becomes indispensable
- **Transparency + ownership mitigates everything**

8.4 Implementation

- Track every implemented recommendation's realized outcome vs predicted outcome
- 2-week look-back loop: flag reversals, scan for negative drift

- Auto-generate "Pryzm self-correction" notes in weekly email
 - Include root-cause hypothesis + next-step investigation plan
-

9. The "Hire an Accountant in India for €500" Objection

The Indian manufacturer meeting raised this — a cost objection Pryzm will hear again.

The wrong answer

"But we have ML!" → nobody cares about capability.

The right answer (expert's framing)

"The answer to this is — this much volume, or this much depth analysis should be done, which accountant cannot do accurately. Which AI can do."

Counter-positioning:

- An accountant in India at €500/month cannot analyze 5,500 invoices × 300 customers × SKU-level margin leakage × cross-customer price discrepancy × seasonality-adjusted forecasting **every night**
- An accountant cannot train 3 forecast models daily and pick the best band per SKU
- An accountant produces a spreadsheet. Pryzm produces **tracked actions with attribution**.
- Pryzm replaces nobody. Pryzm gives the controller and pricing manager they already employ a tool they can't build in Excel.

The geographic point (separate issue)

Expert noted: "If you make in Germany and sell in India, that will always be a problem. Make in India, sell in Germany." → pricing expectation mismatch is structural. Not Pryzm's

problem to solve, but relevant for deal qualification.

10. Phased Rollout — Philosophy Update

Your current phase plan:

- Phase 0: Data cleaning
- Phase 1: Data intelligence
- Phase 2: Forecasting
- Phase 3: Pricing
- Phase 4+: Monte Carlo simulations → "take this action, with this team"

Expert's correction

Don't wait until Phase 5 to show recommendations. Show them early, flagged as early-phase, with transparent error bands.

"Create things, tell them, how do you feel about this? They also know that you are in the early phase. They know that they don't have to trust you blindly. So if they do it, and you do it, then they only know."

Rationale:

- Clients won't leave if an early recommendation is imperfect — they'll leave if they can't see value
- Co-design beats hide-and-reveal
- Feedback from imperfect recommendations trains the next version
- The Indian manufacturer meeting proved this: they gave data, Pryzm showed predictions, some matched, some they'd already exceeded → directional validation

was worth more than waiting for perfection

Tactical implication: Pull Phase 3/4 recommendation surfaces forward into the dashboard NOW, labelled "early preview — help us co-design." Biweekly review cadence with pilot clients (Scherzinger, Avana) as already discussed.

11. The Forecasting Story (How to Pitch the Crown Jewel)

This is what Scherzinger already validated — don't undersell it.

What Pryzm does that nobody competing in this space does

1. **Three years of data → multi-model ensemble** (gradient boosting, time-series FMs, Bayesian for thin-data)
2. **Seasonality detection with web-context enrichment** — model searches web to explain spikes ("Why this spike? What linked event?") and builds a seasonality index (0-1)
3. **Growth trajectory + Pareto repeat-customer modeling** — top 20 customers drive 80% in both Scherzinger and Avana; model weights repeat behavior heavily
4. **Monte Carlo band generation** — conservative / primary / aggressive instead of a single point forecast (this came straight from your physics background — put many dots, connect with linear model, see the envelope)
5. **Walk-forward revalidation** — every new month's data → retrain → predict next quarter → more accurate than static prediction
6. **Model-routing agent (proprietary)** — trains every model every day, picks the winner per SKU per period, produces the final band

The credibility story

- **Avana** (Indian medical distributor, ~30% YoY growth, strong seasonality from hospital annual buying): <2% forecast error on financial-year prediction
- **Scherzinger** (flat revenue precision manufacturer): harder to predict precisely *because* flat — but Q1 2025 conservative/primary/aggressive bands matched actual within 2-3% error
- **Key insight for pitch:** *"Flat-revenue prediction under 5% is harder than seasonal prediction under 2% — we're proving the model works on the hard case first."*

The two-type-of-business caveat (expert agreed)

Not every client will have a clean repeat pattern. Parameter sets will expand per client.

Walk-forward method + latest-data input handles drift. **Don't overclaim universal accuracy.** Claim: "10% error on monthly, under 10% on quarterly — given sufficient data" — that's what the physicist in you can commit to.

12. Margin Leakage — The Phase 2 Use Case That Lands

This is your current-Phase-2 hook. Keep it sharp.

The logic (direct, defensible)

- Customer A bought Article Y at €600 → margin leaking below 25%, unprofitable
- Customer B bought **same** Article Y at same volume at €900
- **Why is Customer A leaking?** → actionable question, data-driven answer, SKU-level

What makes it sellable

- Zero ML black-box objection — it's direct data logic

- SKU-level, customer-level, actionable
- Frames Pryzm as "**cost intelligence / negotiation intelligence platform**" — per the positioning pivot you already identified
- Combines with inventory data: the Indian medical distributor found "hidden" inventory they didn't know they had → unexpected value layer

What to add on top (Phase 2.5)

- **Inventory-adjusted true margin** — if this SKU sat 14 months in warehouse, real margin is lower than invoice margin
 - **Forecast-informed price floor** — "based on next year's forecast, here's the minimum price for this customer"
 - **Cross-customer volume-matched price discrepancy table** — the Excel-competitor killer
-

13. What Corporates Actually Buy (Expert's Three-Job Frame)

Pin this on the wall. This is the entire product thesis in three sentences.

"Your work is three things: 1. Tell the leadership it is a smart tool. What change can it bring in their company? 2. Making it easy for them, bringing this change from their different resources. 3. Sticking this thing on their face again and again — see what we have done for you.

If these three things are done, whatever you sell next will happen with eyes closed."

Job	What it means for Pryzm	What to build
1. Change narrative	MD-view: "Here's the before, here's the after, here's the strategic 12-month promise"	Strategic value proposition on first view; financial value counter (€ captured / € at risk); before/after per persona
2. Make change easy	Action dashboard + assignment + feedback loop	Section 5 — this IS the product
3. Stick it on their face	Weekly auto-email per user + aggregate per manager, biweekly review meeting	Auto-generated email digest; PDF reports (exists); biweekly cadence already started

14. The "Noticed by Mistake" Principle

One of the most underrated insights in the conversation:

"When he comes to the screen by mistake, he will see something which he will remember for a month. And, 'last time, when I came here, I got this idea.' So he will continuously come to your product."

Design rule: Every screen must have a *single surface-level insight* visible without clicking — ideally not generic, ideally specific to that user's business. If the user bounces in for 30 seconds, they leave with something quotable at the next leadership meeting.

Current state: user bounces in, sees 30 KPIs, leaves with nothing quotable.

Tactical fix: "Insight of the week" card per persona, top-right of their view, rotating. Example for Pricing Manager: *"This week's outlier: Article 200832-E sold to Customer 101580 at €4.10 — same article to Customer 102330 at €6.80 last month. Spread: 66%. Investigate?"*

15. Concrete Redesign — Before/After

Landing page — BEFORE

- Dashboard Overview (generic KPIs)
- User scrolls
- User sees Revenue & Margins
- Products & Articles (random-order KPIs)
- Customers
- Forecasting ← the crown jewel, buried here, 5th menu item
- Pricing & Quotes
- ML Analytics
- AI Insights ← "college intern summary"
- Search

Landing page — AFTER (Pricing Manager persona)

PRYZM — Pricing Manager — Week of Apr 20, 2026

STRATEGIC VALUE (next 12 months):

"Margin leakage reduced to <2pp on 80% of SKUs"

Current status: 4.6pp gap on 60% of SKUs

FINANCIAL VALUE (YTD):

€147K captured | €82K at risk | €296K open opportunity

YOUR 3 ACTIONS THIS WEEK (each with 4-option feedback)
[listed, ranked, assignable]

INSIGHT OF THE WEEK (noticed-by-mistake card)
Article 200832-E: 66% price spread same article,
same volume, two customers. [Investigate →]

CUMULATIVE PERFORMANCE (you can forward this to your MD)
Chart: recommendations accepted vs €-captured, 12w

[FORECAST] [QUOTES] [CUSTOMERS] [SEARCH] ← secondary

Same for MD, CFO, Sales Lead, Sales Rep, Inventory, Controller — each persona, each landing, action-first.

16. What to Say to Scherzinger (and Next Clients) — Pitch Reframe

Before (what you say today, rough)

"We have a platform that analyses your margin data, tracks quotes vs invoices, forecasts with <5% error, and generates AI insights across 9 dashboard pages."

→ Client response: "Cool. We have that in Excel."

After (what you should say)

"Pryzm does three things your Excel cannot:

1. **We tell you, each week, the three pricing actions that will most move your margin.**
You decide whether to implement them.
2. **We track what you implemented vs what you didn't** — and show you, in one number, what your team captured and what it left on the table.

3. **Your manager sees this before you send your Monday report** — so Pryzm becomes your team's source of truth.

In the background, we forecast your revenue band to under 5% error, detect SKU-level margin leakage, and flag our own recommendations when they backfire — so you find out from us, not from the P&L two months late."

→ Client response: "Who's the internal owner I should assign?"

That question is the buying signal.

17. Priority Action Items (What to Build Next, Ranked)

P0 — Do this before the next client demo

1. **Action Dashboard + 4-option feedback loop** (Section 5) — this is the product pivot; everything else is polish
2. **Persona router at login** — at minimum MD / Pricing Manager / Sales Lead / Sales Rep distinguished landing pages
3. **Rewrite AI Summary prompt** — "10-year senior pricing manager Monday briefing" pattern per persona
4. **Move forecast to persona-relevant first view** (CFO and Pricing Manager both need the band visible day 1)
5. **Weekly auto-email** summarizing captured-vs-missed per user, aggregated per manager

P1 — Within 4 weeks

6. **Traffic-light sales interface** (Section 7) — plug into existing guardrail feature

7. **Self-correction alerts** (Section 8) — 2-week look-back loop on implemented recommendations
8. **"Insight of the week" card** per persona — the noticed-by-mistake layer
9. **Strategic value proposition headline** per persona landing — define and hard-code three first, iterate with clients
10. **Assignment dropdown** on each action (pre-CRM-integration MVP)

P2 — When CRM/ERP integrations land (SAP SD for Gira, Zoho for Indian client)

11. **Auto-routing of actions** to correct team member per CRM role
12. **Closed-loop impact attribution** — recommendation → implementation → realized € outcome traced end-to-end
13. **Cross-customer volume-matched price discrepancy table** as standalone view
14. **Inventory-adjusted true margin** view (the Indian medical distributor insight)

P3 — Positioning / narrative

15. **Rewrite Pryzm's external positioning** from "margin intelligence platform" → "corporate action layer for pricing, with forecasting and margin leakage detection underneath"
16. **Build the three-jobs narrative** (Section 13) into every investor and client deck
17. **Document the model-routing agent** as proprietary IP — it's your defensible moat, make it legible to sophisticated investors

18. What NOT to Change (The Strong Base)

To keep morale honest — the following are working and should not be touched:

- **Data foundation** — quote-invoice linkage at 89.3%, clean pipeline, quality flags. Keep.
 - **Forecasting band methodology** — Monte Carlo + walk-forward + agent routing is genuinely sophisticated. Keep. Resurface it.
 - **Margin leakage logic** — cross-customer same-volume price discrepancy is a genuine proprietary insight. Keep. Promote it.
 - **Guardrail / live margin alerts feature** — already built, just repurpose as the sales team's primary interface
 - **Biweekly co-design cadence with pilots** — continue, intensify
 - **Phased rollout framework** — keep the structure, pull recommendations forward (Section 10)
 - **Self-hosted open-weight LLM architecture** — don't drift to commercial APIs under pressure; data sovereignty is non-negotiable
-

19. The One-Paragraph Summary to Internalize

Pryzm today is a technically excellent data tool that shows clients what they mostly already know, on screens no one inside their company owns, with an AI summary that narrates instead of synthesizes. The fix isn't more data or more features — it's a pivot to an **action layer**: persona-based first views, a ranked weekly action list with four-option feedback (accepted+implemented / accepted+not / accepted+partial / rejected), cumulative €-captured-vs-€-missed tracking that a user's manager can forward upward, proactive self-correction when recommendations backfire, and a traffic-light simplification for sales. The forecasting and margin-leakage work you've already built is the moat — it just needs to be surfaced as a consequence of action, not as a standalone menu item. The three jobs to internalize: **tell leadership what change this brings, make the change easy to execute, stick it in their face every week.**

20. Quotes to Keep on a Sticky Note

- "What they want to see instead is financial value and strategic value."
- "It should always work backwards from each client."
- "You have thrown everything in their mouth. Who is going to see?"
- "Value-wise there is so much. Structure-wise there is nothing."
- "Apart from the owner, no user works that hard. And the owner is never a user."
- "When he comes to the screen by mistake, he will see something which he will remember for a month."
- "That's how corporate works. I gave you 30 actions. You did 6. Do you know what you lost from the 24?"
- "Your AI summary is now the review of a college intern."
- "Now you're doing fancy Excel. You're doing with capabilities. What do I have to do with your capabilities?"
- "Transparency and ownership mitigates everything."
- "Your work is three things: tell leadership what change it brings; make it easy; stick it on their face again and again."

End of strategic playbook. Next step: review Section 17 (Priority Action Items) and identify which PO items can ship before the next Scherzinger biweekly.