

Scherzinger Margin Intelligence Platform

Complete Project Deep Dive

Prepared for: Shahad (Client Demo Preparation)
Date: April 1, 2026
Status: Phases 1–3 Complete | Frontend + Backend Remaining Work | Demo Ready

Contents

1	What Is This Project?	2
1.1	Tech Stack	2
2	Where Did the Data Come From?	2
3	What Was in the Raw Data?	2
3.1	Invoice Data (Deckungsbeitragsliste)	2
3.1.1	What DB I and DB II Mean (German Accounting)	3
3.2	Quote Data (Angebotsstatistik)	3
3.3	Rejection Codes (15 Codes)	4
3.4	The 9 Commodity Groups	5
4	How We Cleaned & Prepared the Data	5
4.1	The Cleaning Pipeline (<code>clean_data.py</code>)	5
4.2	Data Loading (<code>load_data.py</code>)	6
4.3	Quote-Invoice Linkage (<code>link_quotes_invoices.py</code>)	6
5	Database Schema—What We Built	7
5.1	Phase 1 Tables (6)	7
5.2	Phase 2 Tables (7)	7
6	Phase 1—Core Margin Analytics	7
6.1	What Questions Can the Client Answer with Phase 1?	7
6.2	Key Phase 1 Findings (Verified Numbers)	8
7	Phase 2—Forecasting, Risk & ML Models	9
7.1	Margin Forecasting (4 Models)	9
7.1.1	Model 1: Exponential Moving Average (EMA)	9
7.1.2	Model 2: Linear Trend	9
7.1.3	Model 3: Seasonal Decomposition	9
7.1.4	Model 4: Ensemble (Weighted Average)	9
7.2	Customer Risk Scoring (5 Components)	9
7.3	Monte Carlo Simulations	10
7.4	Seasonal Pattern Analysis	11
7.5	Commodity Benchmarks	11
7.6	Cost Trend Analysis	11
7.7	Walk-Forward Backtests	11
8	Phase 3—Frontend Dashboard	12
8.1	Pages Built	12
9	Every Graph & Plot Explained	12
9.1	Dashboard KPI Cards (4 cards at top)	12
9.2	Monthly Revenue & Margin Trend (Line + Area Chart)	12
9.3	Revenue by Customer (Horizontal Bar Chart)	12
9.4	Revenue by Product (Horizontal Bar Chart)	12
9.5	Commodity Group Performance (Bar + Pie Chart)	12
9.6	Win Rate by Deal Size (Bar Chart)	13
9.7	Rejection Code Distribution (Horizontal Bar)	13
9.8	Risk Tier Distribution (Pie / Stacked Bar)	13
9.9	Forecast Model Comparison (Multi-Line Chart)	13

9.10	Forecast Confidence Bands (Area Chart)	13
9.11	Monte Carlo Distribution (Histogram / Box Plot)	13
9.12	Seasonal Pattern Heatmap	13
9.13	Cost Trend Lines (Line Chart per Product)	13
9.14	Backtest Accuracy Comparison (Bar Chart)	13
10	All Formulas & Calculations	13
10.1	Core Margin Formulas (German Accounting)	13
10.2	Cost Component Breakdown	14
10.3	Quote–Invoice Margin Gap	14
10.4	Revenue-Weighted Margin	14
10.5	Win Rate	14
10.6	Price Sensitivity T-Test	14
10.7	EMA Forecast Formula	14
10.8	Linear Trend Forecast Formula	14
10.9	Seasonal Decomposition Forecast	15
10.10	Ensemble Forecast	15
10.11	Risk Score Components	15
10.12	Monte Carlo Simulation	16
10.13	Seasonal Index	16
10.14	Cost Change Percentage	16
10.15	Backtest Metrics	16
11	ML Models—How They Work, Accuracy & How We Got There	17
11.1	Why These Models?	17
11.2	Model Development Process	17
11.3	Accuracy Summary	17
11.4	Current Forecast Values	17
12	JSON Data Files—What They Contain	18
13	Example SKU Scenario 1: Article 300143 (Declining Margin, At Risk)	18
13.1	The Product	18
13.2	Year-by-Year Performance	19
13.3	What Happened?	19
13.4	Cost Breakdown	19
13.5	Why It's At Risk	19
13.6	What the Platform Tells Scherzinger to Do	19
13.7	Revenue Impact if Nothing Changes	19
14	Example SKU Scenario 2: Article 201885 (Healthy & Growing)	20
14.1	The Product	20
14.2	Year-by-Year Performance	20
14.3	What's Going Right?	20
14.4	Cost Breakdown	20
14.5	Why It's Healthy	20
14.6	Revenue Potential	20
14.7	Comparison: SKU 300143 vs 201885	21
15	How Much Does This Save Them?	21
15.1	Direct Savings from Margin Visibility	21
15.2	Total Estimated Annual Impact (Phases 1–3)	21
15.3	Time Savings	21

16 Future Phases—What They Can Expect	22
16.1 Phase 4: Pricing & FX Analysis (Planned)	22
16.2 Phase 5: Advanced ML Analytics (Planned)	22
16.3 Phase 6+ (Potential Future)	22
17 Future Phase Example Scenarios	22
17.1 Scenario A: Phase 4—Price Optimization for Article 300143	22
17.2 Scenario B: Phase 5—ML Churn Prediction for Customer 101690	23
17.3 Scenario C: Phase 5—Demand Classification for Production Planning	23
18 API Endpoint Reference	24
18.1 Phase 1 Endpoints (19)	24
18.2 Phase 2 Endpoints (9)	24
19 Key Numbers to Remember for the Demo	25
19.1 The Big Numbers	25
19.2 The Problem Numbers (What They Need to Fix)	25
19.3 The Opportunity Numbers (What They Can Gain)	25
19.4 The Forecast Numbers	25
19.5 The Model Numbers	25

Section 1: What Is This Project?

Scherzinger GmbH is an industrial pump manufacturer. They have 4 years (2022–2025) of historical invoice and quotation data. Before this platform, all margin analysis was manual—spreadsheets, no forecasting, no risk awareness, no way to see which customers or products were silently losing them money.

We built a **Margin Intelligence Platform** that:

- **Phase 1:** Ingests their raw Excel data, cleans it, loads it into a PostgreSQL database, and provides 19 API endpoints for margin analysis, quote performance, and data quality reporting.
- **Phase 2:** Adds predictive intelligence—4 forecasting models, a 5-component customer risk scoring system, Monte Carlo margin simulations, seasonal pattern detection, cost trend tracking, commodity benchmarking, and walk-forward backtest validation.
- **Phase 3:** Builds a React frontend dashboard with 9 pages that visualize all of Phase 1 and 2’s analytics—KPI cards, trend charts, customer/product tables, forecast comparison views, risk tier breakdowns, and more.

1.1 Tech Stack

- **Backend:** Python (FastAPI 0.109), PostgreSQL 15, SQLAlchemy 2.0
- **Frontend:** React 19.2 + Vite 7, Recharts 3.7, Tailwind CSS 4.2, Framer Motion
- **ML / Analytics:** NumPy, Pandas, SciPy, Statsmodels
- **Data Pipeline:** Openpyxl (Excel reading), PyArrow (Parquet), Alembic (DB migrations)

Section 2: Where Did the Data Come From?

Scherzinger provided **3 Excel files**:

File	What It Contains	Records	Sheets
Deckungsbeitragsliste_2.xlsx	Invoice / contribution margin data	5,565	4 (2022–2025)
Angebotsstatistik_3.xlsx	Quotation statistics (won/lost)	4,605	4 (2022–2025)
Quotation code interpretation.xlsx	Rejection code reference table	15	1

Table 1: Source data files provided by Scherzinger.

All data is in **German** (column names, rejection codes, product descriptions). We mapped everything to English for the platform.

Section 3: What Was in the Raw Data?

3.1 Invoice Data (Deckungsbeitragsliste)

Each row = one line item on an invoice. Here is what Scherzinger tracks per invoice line:

German Column	English Meaning	Example	What It Tells Us
Rechnung	Invoice number	6009811	Unique invoice ID
Pos.	Line item position	1, 2, 3	Multiple items per invoice
Datum	Invoice date	2022-01-15	When the sale happened
Firma	Customer ID	103459	Which customer bought
WG	Commodity group	BKAES	Product category (9 groups)
Artikel	Article/SKU number	202427	Specific product
Menge	Quantity	1–1000+	Units sold
Umsatz	Revenue (EUR)	10K–500K	Total sales value
HKvoll / Stck.	Full mfg cost/unit	100–5,000	Total cost to make one unit
HKvar / Stck.	Variable mfg cost/unit	50–3,000	Variable portion of cost
MatAnteil / Stck.	Material cost/unit	20–2,000	Raw materials cost
FEK / Stck.	Direct mfg cost/unit	30–1,000	Direct labor + machine cost
davon FV / Stck.	Outsourcing cost/unit	5–500	External production cost
DB I	Contribution Margin I	5K–250K	Revenue – variable costs
DB II	Contribution Margin II	2K–200K	Revenue – ALL costs (key metric)
DB I Marge	DB I as %	0.20–0.90	Variable margin ratio
DB II Marge	DB II as %	–0.19 to 0.92	THE PRIMARY METRIC

3.1.1 What DB I and DB II Mean (German Accounting)

- **DB I (Deckungsbeitrag I)** = Revenue – Variable Costs. This tells you: “after paying for materials and direct labor, how much is left?”
- **DB II (Deckungsbeitrag II)** = Revenue – ALL Costs (variable + fixed overhead). This is the **real profitability**—it accounts for factory overhead, salaries, depreciation, everything. **This is the metric the entire platform is built around.**

3.2 Quote Data (Angebotsstatistik)

Each row = one line item on a quotation sent to a customer.

German Column	Col-	English Meaning	Example	What It Tells Us
Angebot		Quote ID	AN102237	Unique quote identifier
Pos		Position	1, 2	Line item
Datum		Quote date	2022-03-20	When quote was sent
Firma		Customer ID	103459	Who we quoted
Artikel		Article/SKU	202427	What product
Menge		Quantity	1–1000+	How many they wanted
Umsatz		Quoted revenue	15K–600K	What we quoted them
HKvoll		Total mfg cost	5K–250K	Our cost (NOT per-unit!)
DB2%		DB II margin %	73.26	Stored as % (73.26 = 73.26%)
Status Code		4=Won, 5=Lost	4 or 5	Did we win the deal?
Auftrag		Order ID	3013751	Links to invoice (if won)

German um	Col-	English Meaning	Example	What It Tells Us
Rejection Code		Loss reason	PA, PR, KA	Why we lost (if lost)

Critical note: Quote DB2% is stored as a raw percentage (e.g., 73.26), while invoice DB2 Marge is stored as a decimal (e.g., 0.7326). Our cleaning script divides quote values by 100 to normalize.

3.3 Rejection Codes (15 Codes)

These explain WHY a quote was lost. Only reliable from 2025 onward (earlier years have incomplete codes).

Code	German	English	Pricing?	Model Use
PA	Parallelangebot	Competitor cheaper	YES	Risk scoring
PR	Preis	Price too high	YES	Risk scoring
KE	Keine Reaktion Endkunde	End cust. no response	Cautious	Weighted lower
KN	Kundenprojekt n. realisiert	Project cancelled	Cautious	External factor
KR	Keine Reaktion	No response	Cautious	Unknown cause
QS	Qualität	Quality concerns	Cautious	Perception
AN	Anfrage	Inquiry only	Cautious	Low intent
FI	Firmenimage	Company image	Cautious	Perception
LZ	Lieferzeit	Delivery too long	Exclude	Supply chain
DO	Dokumentation	Missing docs	Exclude	Regulatory
TE	Techn. nicht passend	Technical mismatch	Exclude	Product fit
RZ	Reaktionszeit	Slow response	Exclude	Internal
SL	Systemlieferant	System supplier	Exclude	Relationship
KA	Keine Angabe	No information	Exclude	Unknown
KD	Kunde n. kontaktiert	No follow-up	Exclude	Internal

Why this matters: Codes PA and PR tell us “we lost because of price.” These feed into the risk scoring model’s rejection component.

3.4 The 9 Commodity Groups

Group	Full Name	Revenue (4yr)	Invoices	Avg DB2	BCG
BKAES	Gear pumps (Zahnrad- pumpen)	€ 16,706,796	2,819	66.8%	Star (+2.7%)
BKAGG	Internal gear ring pumps	€ 6,811,773	2,545	53.8%	Dog (−9.6%)
BKAIZ	Industrial gears	€ 811,613	159	56.1%	Question Mark (+70.9%)
SOPU	Special pumps	€ 170,000	12	46.4%	Dog (−96.2%!)
MBDIV	Misc. com- ponents	€ 107,024	14	59.1%	Question Mark (+21%)
SOPUZK	Specialty units	€ 16,524	8	81.0%	Cash Cow
OFRSCR	Frames/structures	€ 10,236	2	66.1%	Cash Cow
MBKUEHL	Cooling components	€ 4,062	2	9.8%	Dog (low margin!)
OFRLMG	Logistics/handling	€ 693	1	64.5%	Cash Cow

Table 5: Commodity groups with revenue, margin, and BCG classification.

Key insight for demo: BKAES is their bread and butter (68% of all revenue, healthy margins). BKAGG is the second-largest but declining. BKAIZ is small but growing fast—a potential future star. SOPU has essentially collapsed (−96% growth).

Section 4: How We Cleaned & Prepared the Data

4.1 The Cleaning Pipeline (`clean_data.py`)

Step 1—Read Excel Sheets: We read all 4 yearly sheets from both Excel files (2022, 2023, 2024, 2025) and concatenate them.

Step 2—Column Mapping: All German column names mapped to English. Types converted: dates parsed, numbers coerced, strings stripped.

Step 3—Derived Fields Added:

- `year`, `quarter`, `month` extracted from the date field
- `business_unit` set to “BU001” (single business unit)
- Quote `db2_margin` = `db2_pct_raw` / 100.0
- Quote `is_won` = True if `status_code` == 4, False if 5

Step 4—Data Quality Flags:

For invoices:

```

dq_missing_margin = TRUE if db2_margin is NULL      (20 records ,
    0.36%)
dq_negative_margin = TRUE if db2_margin < 0          (13 records ,
    0.23%)
dq_low_margin      = TRUE if 0 <= db2_margin < 0.10 (flagged but not
    excluded)
dq_any_issue       = TRUE if any of the above is TRUE

```


For quotes:

```
dq_100pct_margin = TRUE if db2_pct_raw == 100.0 (802 records, 17.4%)
  -> These are quotes where HKvoll = 0, so DB2 = revenue, giving 100%
      "margin"
  -> They are EXCLUDED from all model training
dq_missing_cost  = TRUE if hkvoll is NULL or 0 (96 records)
dq_any_issue     = TRUE if any of the above is TRUE
```

Step 5—Rejection Code Normalization:

- Mixed case codes normalized to uppercase (“ka” → “KA”, “Pa” → “PA”)
- Short codes mapped: “P” → “PR”, “T” → “TE”
- Reliability flag: `rejection_code_reliable` = TRUE only for 2025 data

Step 6—Deduplication: Quotes deduplicated on (`quote_id`, `position`), keeping the first occurrence.

Step 7—Output: Clean data written to Parquet files:

- `invoices_clean.parquet` (5,565 records)
- `quotes_clean.parquet` (4,539 records after dedup)
- `customers.parquet` (967 unique)
- `products.parquet` (1,223 unique)

4.2 Data Loading (`load_data.py`)

Parquet files loaded into PostgreSQL via batch inserts (1,000 records per batch). Tables created via SQLAlchemy ORM + Alembic migrations.

4.3 Quote–Invoice Linkage (`link_quotes_invoices.py`)

This is one of the most valuable data transformations. We link won quotes to their resulting invoices using the `order_id` (Auftrag) field:

How it works:

1. Take all won quotes that have an `order_id`
2. Find all invoices with matching `order_id`
3. Create a link record capturing: quoted margin, actual margin, the gap, and days between quote and invoice

Linkage Results:

- Won quotes with `order_id`: 1,724 (99.5% of won quotes)
- Successfully matched to invoices: 1,406 (89.3%)
- Unmatched: 169 (53 were service orders with “S-” prefix, 116 were recent/pending)
- Total link records created: 1,957

Margin Gap Calculation:

```
margin_gap = quoted_db2_margin - actual_db2_margin
```

- **Positive gap** = we quoted higher than what we actually invoiced (margin erosion)
- **Negative gap** = actual margin was better than quoted (favorable)

Key Findings:

- Mean margin gap: **+1.9pp** (we consistently quote slightly higher than we deliver)
- Gap by year: 2022: 2.3pp, 2023: 1.9pp, 2024: 2.1pp, 2025: 1.4pp (improving!)
- Average days from quote to invoice: **~62 days** (median ~46 days)

Section 5: Database Schema—What We Built

14 tables total across 2 phases:

5.1 Phase 1 Tables (6)

1. **customers**—967 records: `customer_id`, `first_seen_date`
2. **products**—1,223 records: `article_id`, `description`, `commodity_group`, `drawing`
3. **invoices**—5,565 records: 26 columns covering revenue, costs (5 cost components), margins (DB1, DB2), quantities, dates, quality flags
4. **quotes**—4,539 records: 19 columns covering quoted values, win/loss status, rejection codes, quality flags
5. **quote_invoice_links**—1,957 records: the bridge between quotes and invoices with `margin_gap`, `days_to_invoice`
6. **rejection_codes**—15 records: `code`, German name, English name, pricing impact classification

5.2 Phase 2 Tables (7)

7. **margin_forecasts**—~1,500 records: predictions from 4 models at 1/3/6/12 month horizons
8. **customer_risk_scores**—~180 records: composite 0–1 risk score with 5 component breakdowns
9. **product_cost_trends**—~700 records: quarterly cost evolution per product
10. **seasonal_patterns**—~300 records: monthly seasonal indices
11. **commodity_benchmarks**—~80 records: benchmark margins per commodity group per year/quarter
12. **monte_carlo_results**—~130 records: stochastic simulation results with percentile distributions
13. **backtest_results**—~40 records: walk-forward validation metrics by model type

Section 6: Phase 1—Core Margin Analytics

Phase 1 delivers **19 API endpoints** that answer fundamental business questions:

6.1 What Questions Can the Client Answer with Phase 1?

Margin Analysis:

- “What’s our overall margin?” → `/margins/summary`
- “How has our margin changed year over year?” → `/margins/by-year`
- “Who are our most profitable customers?” → `/margins/by-customer`
- “Which products make us the most money?” → `/margins/by-product`
- “How do our 9 commodity groups compare?” → `/margins/by-commodity-group`
- “How much margin do we lose between quoting and invoicing?” → `/margins/gap-analysis`
- “What’s the trend—are we improving?” → `/margins/trend`

Quote Performance:

- “What’s our win rate? Improving?” → /quotes/summary + /quotes/win-rate-by-year
- “Do we win more on small or large deals?” → /quotes/win-rate-by-deal-size
- “Why are we losing deals?” → /quotes/rejection-codes
- “Are we pricing ourselves out?” → /quotes/price-sensitivity

6.2 Key Phase 1 Findings (Verified Numbers)

Metric	Value
Total 4-year revenue	€ 24,646,718
Total invoices	5,565
Total quotes	4,605 (4,539 after dedup)
Overall DB2 margin (mean)	64.78%
Overall win rate	37.6% (1,733 won / 4,605 total)
Unique customers	967
Unique products / SKUs	1,223
Commodity groups	9
Quote–invoice match rate	89.3%
Mean margin gap	+1.9pp

Table 6: Phase 1 key metrics.

Win Rates by Year (improving trend):

- 2022: 36.5% (346 won / 947 total)
- 2023: 34.3% (408 won / 1,191 total)
- 2024: 38.9% (458 won / 1,176 total)
- 2025: **40.4%** (521 won / 1,291 total)

Revenue by Year:

- 2022: €6,369,103 (1,500 invoices, 448 customers, avg margin 63.6%)
- 2023: €6,233,961 (1,337 invoices, 440 customers, avg margin 63.8%)
- 2024: €5,793,294 (1,320 invoices, 391 customers, avg margin 62.2%)—dip year
- 2025: €6,250,360 (1,408 invoices, 411 customers, avg margin 61.0%)—recovery

Lost Revenue by Rejection Code (2025 data, reliable):

- KA (No info): €4,942,412—51.7% of all losses
- KR (No response): €1,318,560—18.4%
- AN (Inquiry only): €1,168,322—9.2%
- **PA (Competitor cheaper): €793,893—5.8%** ← Price competition
- KE (End customer silent): €274,121—5.0%
- KD (We didn’t follow up): €369,629—4.0% ← Internal failure!
- KN (Project cancelled): €503,198—2.9%
- **PR (Our price too high): €177,374—1.8%**
- TE (Technical mismatch): €22,195—0.6%
- LZ (Delivery too slow): €14,627—0.3%

Price Sensitivity Analysis (t-test):

- Won quotes average margin: 70.6%
- Lost quotes average margin: 72.4%
- Difference: 1.8pp (lost quotes were priced higher)
- P-value: **0.006** (statistically significant!)

- Meaning: There IS a measurable relationship between higher pricing and losing deals.

Section 7: Phase 2—Forecasting, Risk & ML Models

Phase 2 adds **predictive intelligence** via 9 additional API endpoints and 7 new database tables.

7.1 Margin Forecasting (4 Models)

We forecast DB2 margin forward at 1, 3, 6, and 12 month horizons for: Top 200 customers (by revenue), Top 300 products (by revenue), and all 9 commodity groups.

7.1.1 Model 1: Exponential Moving Average (EMA)

- Requires: ≥ 12 months of data
- How: Applies exponential weighted mean with span=6, giving more weight to recent months
- Confidence interval: $\pm 1.645 \times \sigma_{\text{residuals}}$ (90% CI)
- Good for: Short-term smoothing, reacting to recent changes

7.1.2 Model 2: Linear Trend

- Requires: ≥ 18 months of data
- How: Fits a straight-line regression (OLS) on monthly margin values over time
- Formula: $\hat{m} = \beta_0 + \beta_1 \times t$
- Reports R^2 (goodness of fit)
- Good for: Detecting directional trends (is margin going up or down?)

7.1.3 Model 3: Seasonal Decomposition

- Requires: ≥ 24 months of data
- How: Uses statsmodels additive seasonal decomposition with 12-month period
- Separates margin into: trend + seasonal + residual
- Prediction = extrapolated trend + seasonal factor for target month
- Good for: Capturing repeating monthly patterns

7.1.4 Model 4: Ensemble (Weighted Average)

- Requires: ≥ 18 months of data
- If seasonal available: **30% EMA + 30% Linear + 40% Seasonal**
- If seasonal unavailable: **50% EMA + 50% Linear**
- Confidence interval: [min(all lowers), max(all uppers)]
- Good for: Robust predictions that balance multiple perspectives

All predictions clipped to $[-1.0, 1.0]$ range.

7.2 Customer Risk Scoring (5 Components)

Every qualifying customer (≥ 5 invoices, ≥ 3 quote–invoice links, ≥ 5 quotes) gets a **composite risk score from 0.0 to 1.0**.

Component	Weight	What It Measures	Formula
Margin Trend	30%	Is margin going up or down?	$\text{clip}\left(\frac{0.05-\Delta}{0.15}, 0, 1\right)$
Margin Gap	25%	Quote-to-invoice erosion	$\text{clip}\left(\frac{\bar{g}}{0.15}, 0, 1\right)$
Volume Trend	20%	Purchase volume growth/decline	$\text{clip}\left(\frac{0.10-\Delta_{\text{rev}}}{0.40}, 0, 1\right)$
Win Rate	15%	Are we winning their quotes?	$\text{clip}\left(\frac{0.50-w_r}{0.35}, 0, 1\right)$
Rejection Pattern	10%	Price-driven losses (PA+PR)	$\text{clip}\left(\frac{r_p-0.10}{0.40}, 0, 1\right)$

Table 7: Risk scoring components and formulas.

Risk Tiers:

Tier	Score Range	Count	Avg Score
Low	0.00–0.24	33 customers (18%)	0.22
Medium	0.25–0.49	63 customers (34.5%)	0.52
High	0.50–0.74	65 customers (35.6%)	0.74
Critical	0.75–1.00	22 customers (12%)	0.91

Table 8: Customer risk tier distribution.

What this means: 47.6% of scored customers are at High or Critical risk.

7.3 Monte Carlo Simulations

We run stochastic simulations to model the probability distribution of future margins.

How it works (step by step):

1. Take a customer/product/group’s monthly margin history
2. Compute monthly returns: $r_t = m_t - m_{t-1}$
3. Calculate drift ($\mu = \bar{r}$) and volatility ($\sigma = \text{std}(r)$)
4. Simulate 10,000 paths (5,000 for customers): each month $m_{t+1} = \text{clip}(m_t + \mathcal{N}(\mu, \sigma), -1, 1)$
5. From 10,000 final margins, compute percentiles ($p_5, p_{25}, p_{50}, p_{75}, p_{95}$)
6. Compute $\text{Pr}(\text{margin} < 50\%)$

Horizons: 3, 6, and 12 months. **Random seed:** 42.

Entity	Mean	Median	P5	P95	Pr < 50%
Overall	59.5%	59.4%	35.9%	83.0%	25.6%
BKAES	64.9%	64.8%	39.4%	90.3%	16.9%
BKAGG	43.9%	43.7%	7.0%	80.7%	61.0%
BKAIZ	53.1%	55.8%	−6.1%	100%	43.4%
MBDIV	0.1%	3.1%	−100%	100%	52.7%
SOPU	2.2%	5.0%	−100%	100%	72.5%
SOPUZK	56.5%	64.8%	−18.8%	100%	37.3%

Table 9: Monte Carlo simulation results (12-month horizon).

Key insight: For BKAGG (second-largest group), there is a **61% probability** that margins will fall below 50% within 12 months.

7.4 Seasonal Pattern Analysis

Seasonal index for each month: $I_m = \bar{m}_{\text{month}} / \bar{m}_{\text{grand}}$.

Three levels: (1) Overall, (2) By commodity group, (3) By qualifying customer (≥ 48 invoices).

Overall seasonal range: 0.962 (weakest) to 1.037 (strongest).

7.5 Commodity Benchmarks

For each commodity group, per year and quarter, we calculate: simple average DB2 margin, revenue-weighted DB2 margin, median/P25/P75 percentile margins, total revenue, record count, average win rate, and average margin gap.

7.6 Cost Trend Analysis

For products with ≥ 3 quarters of data and ≥ 2 invoices per quarter:

- Track quarterly quantity-weighted averages: HKvoll, material, FEK, FV costs per unit
- Calculate: $\Delta_{\text{cost}} = (C_{\text{current}} - C_{\text{prior}}) / |C_{\text{prior}}|$
- Flag warning if $|\Delta_{\text{cost}}| > 50\%$

7.7 Walk-Forward Backtests

1. Train on first N months (min 12/18/24 depending on model)
2. Predict the next 3-month average
3. Expand window by 3 months, repeat
4. Compare predictions to actuals

Model	MAE	RMSE	Dir. Accuracy
Linear Trend	0.024	0.028	54.5%
Seasonal Decomp	0.032	0.037	9.1%
Ensemble	0.035	0.040	9.1%
EMA	0.065	0.068	54.5%

Table 10: Backtest accuracy results.

MAE of 0.024 means predictions are off by $\sim 2.4\text{pp}$ on average. **Directional accuracy of 54.5%** means correct up/down prediction $\sim 55\%$ of the time.

Section 8: Phase 3—Frontend Dashboard

8.1 Pages Built

Page	Route	What It Shows
Dashboard	/	KPI cards, monthly trend, top customers, commodity breakdown
Revenue & Margins	/revenue-margins	Detailed margin trends, YoY comparisons, commodity performance
Customers	/customers	Customer table with revenue, margin, win rate, risk tier
Products	/products	SKU-level performance, cost breakdowns, at-risk flags
Forecasting	/forecasting	Model comparison, forecast bands with 90% CI, Monte Carlo
Quality Metrics	/quality	Data quality dashboard, missing data, anomaly detection
AI Insights	/ai-insights	BCG matrix, margin classification, customer risk predictions
Pricing & FX	/pricing	<i>Stub for Phase 4</i>
ML Analytics	/ml-analytics	<i>Stub for Phase 5</i>

Table 11: Frontend pages.

All charts built with **Recharts 3.7**. Animations via **Framer Motion**. Responsive via **Tailwind CSS 4.2**.

Section 9: Every Graph & Plot Explained

9.1 Dashboard KPI Cards (4 cards at top)

- **Total Revenue:** Sum of all invoices. Currently €24.6M (all time) or €6.25M (2025).
- **Average DB2 Margin:** Mean margin. Currently 64.78% all-time, 61.0% in 2025.
- **Win Rate:** Won/total quotes. Currently 37.1%.
- **At-Risk Customers:** High + Critical tier. Currently 87 (65 high + 22 critical).

9.2 Monthly Revenue & Margin Trend (Line + Area Chart)

- X-axis: Month (Jan 2022 → Dec 2025, 48 points)
- Y-axis Left: Revenue in EUR (bars/area); Y-axis Right: DB2 margin % (line)
- Data source: `monthly_detail.json`
- Key: Revenue dipped mid-2024, recovered. Margin slight downward trend 63.6% → 61.0%.

9.3 Revenue by Customer (Horizontal Bar Chart)

Top 20 customers ranked by total revenue. Customer 101690 leads at €1.54M but only 53.3% margin.

9.4 Revenue by Product (Horizontal Bar Chart)

Top 20 SKUs ranked by total revenue. Product 201924-F leads at €580K with 71.7% margin. Product 206028-01 has €410K but only 9.8% margin.

9.5 Commodity Group Performance (Bar + Pie Chart)

Revenue share and margin across 9 groups. Calculation: $\text{weighted margin} = \sum \text{db2_total} / \sum \text{revenue}$. BKAES dominates (68% revenue, 66.8% margin).

9.6 Win Rate by Deal Size (Bar Chart)

Deal size bands: <€1K, €1–5K, €5–10K, €10–50K, >€50K. Sweet spot is €5–10K range.

9.7 Rejection Code Distribution (Horizontal Bar)

Count and revenue impact of each loss reason. KA (51.7%) and KR (18.4%) dominate. PA+PR = 7.6% but €971K.

9.8 Risk Tier Distribution (Pie / Stacked Bar)

Customer distribution across Low/Medium/High/Critical. Only 18% are Low risk.

9.9 Forecast Model Comparison (Multi-Line Chart)

EMA, Linear, Seasonal, Ensemble predictions side by side with 90% CI shaded bands.

9.10 Forecast Confidence Bands (Area Chart)

Central prediction with P5–P95 shaded area. Wider bands = more uncertainty.

9.11 Monte Carlo Distribution (Histogram / Box Plot)

Distribution of 10,000 simulated outcomes. Vertical line at 50% threshold. Shaded area = probability below threshold.

9.12 Seasonal Pattern Heatmap

Month × Commodity group. Color: seasonal index (green > 1.0, red < 1.0).

9.13 Cost Trend Lines (Line Chart per Product)

Quarterly HKvoll per unit with material, FEK, FV breakout lines.

9.14 Backtest Accuracy Comparison (Bar Chart)

MAE and RMSE bars for each model type. Linear Trend wins with MAE 0.024.

Section 10: All Formulas & Calculations

10.1 Core Margin Formulas (German Accounting)

$$\text{DB I} = \text{Revenue} - (\text{HKvar_per_unit} \times \text{Quantity}) \quad (1)$$

$$\text{DB II} = \text{Revenue} - (\text{HKvoll_per_unit} \times \text{Quantity}) \quad (2)$$

$$\text{DB I Margin} = \frac{\text{DB I Total}}{\text{Revenue}} \quad (3)$$

$$\text{DB II Margin} = \frac{\text{DB II Total}}{\text{Revenue}} \quad \leftarrow \text{THE KEY METRIC} \quad (4)$$

10.2 Cost Component Breakdown

$$\text{Material \%} = \frac{\text{material_per_unit}}{\text{hkvoll_per_unit}} \times 100 \quad (5)$$

$$\text{Direct Mfg \%} = \frac{\text{fek_per_unit}}{\text{hkvoll_per_unit}} \times 100 \quad (6)$$

$$\text{Outsourcing \%} = \frac{\text{fv_per_unit}}{\text{hkvoll_per_unit}} \times 100 \quad (7)$$

10.3 Quote-Invoice Margin Gap

$$\text{margin_gap} = \text{quoted_db2_margin} - \text{actual_db2_margin} \quad (8)$$

$$\text{days_to_invoice} = \text{invoice_date} - \text{quote_date} \text{ (in days)} \quad (9)$$

10.4 Revenue-Weighted Margin

$$\text{weighted_margin} = \frac{\sum \text{db2_total}}{\sum \text{revenue}}$$

This is more accurate than simple averaging because a €500K invoice at 70% counts much more than a €500 invoice at 90%.

10.5 Win Rate

$$\text{win_rate} = \frac{\text{COUNT}(\text{is_won} = \text{TRUE})}{\text{COUNT}(\text{all quotes})}$$

10.6 Price Sensitivity T-Test

```
Group A: margins of all WON quotes (excluding dq_100pct_margin)
Group B: margins of all LOST quotes (excluding dq_100pct_margin)
Test: Independent samples t-test (scipy.stats.ttest_ind)
Result: p-value = 0.006 (significant at alpha=0.05)
Interpretation: Lost quotes had higher margins -> pricing affects win
probability
```

10.7 EMA Forecast Formula

```
ema = series.ewm(span=6, adjust=False).mean()
predicted = ema.iloc[-1] # last EMA value projected forward
residuals = series - ema_fitted
ci_width = 1.645 * residuals.std() # 90% confidence
lower = predicted - ci_width
upper = predicted + ci_width
```

10.8 Linear Trend Forecast Formula

```
from scipy.stats import linregress
x = range(len(series))
slope, intercept, r_value, p_value, std_err = linregress(x, series)
```

```

r_squared = r_value ** 2
predicted = intercept + slope * (len(series) + months_ahead - 1)
ci_width = 1.645 * std(residuals)

```

10.9 Seasonal Decomposition Forecast

```

from statsmodels.tsa.seasonal import seasonal_decompose
result = seasonal_decompose(series, model='additive', period=12)
# Extrapolate trend via linear regression on trend component
# Add seasonal factor for the target month
predicted = extrapolated_trend + seasonal_component[target_month]

```

10.10 Ensemble Forecast

```

# With seasonal (requires >= 24 months):
weights = [0.3, 0.3, 0.4] # EMA, Linear, Seasonal
predicted = 0.3 * ema_pred + 0.3 * linear_pred + 0.4 * seasonal_pred

# Without seasonal (requires >= 18 months):
weights = [0.5, 0.5] # EMA, Linear
predicted = 0.5 * ema_pred + 0.5 * linear_pred

# Confidence interval spans the widest of all component intervals
lower = min(ema_lower, linear_lower, seasonal_lower)
upper = max(ema_upper, linear_upper, seasonal_upper)

```

10.11 Risk Score Components

```

# Component 1: Margin Trend (30%)
delta = mean(last 4 quarters) - mean(first 4 quarters)
margin_trend_score = clip((0.05 - delta) / 0.15, 0, 1)

# Component 2: Margin Gap (25%)
gap_score = clip(avg_gap / 0.15, 0, 1)

# Component 3: Volume Trend (20%)
change = (recent_revenue - prior_avg) / prior_avg
volume_score = clip((0.10 - change) / 0.40, 0, 1)

# Component 4: Win Rate (15%)
winrate_score = clip((0.50 - recent_win_rate) / 0.35, 0, 1)

# Component 5: Rejection Pattern (10%)
price_ratio = count(PA + PR) / count(all lost)
rejection_score = clip((price_ratio - 0.10) / 0.40, 0, 1)

# Composite Score
composite = 0.30*margin_trend + 0.25*gap + 0.20*volume
            + 0.15*winrate + 0.10*rejection

```

The composite score formula in mathematical notation:

$$S = 0.30 S_{\text{margin}} + 0.25 S_{\text{gap}} + 0.20 S_{\text{volume}} + 0.15 S_{\text{winrate}} + 0.10 S_{\text{rejection}}$$

10.12 Monte Carlo Simulation

```

returns = np.diff(margin_series) # monthly changes
mu = returns.mean()             # drift
sigma = returns.std()           # volatility
np.random.seed(42)

for sim in range(10000):
    margin = last_known_margin
    for month in range(horizon):
        shock = np.random.normal(mu, sigma)
        margin = clip(margin + shock, -1.0, 1.0)
    final_margins.append(margin)

# Results
p5  = np.percentile(final_margins, 5)
p25 = np.percentile(final_margins, 25)
p50 = np.median(final_margins)
p75 = np.percentile(final_margins, 75)
p95 = np.percentile(final_margins, 95)
prob_below_threshold = np.mean(final_margins < 0.50)

```

10.13 Seasonal Index

$$I_m = \frac{\bar{m}_{\text{month}}}{\bar{m}_{\text{grand}}}$$

If sample count < minimum threshold, index defaults to 1.0.

10.14 Cost Change Percentage

$$\bar{C} = \frac{\sum(\text{hkvoll_per_unit} \times \text{quantity})}{\sum \text{quantity}} \quad \Delta_{\text{cost}} = \frac{\bar{C}_{\text{current}} - \bar{C}_{\text{prior}}}{|\bar{C}_{\text{prior}}|}$$

Warning flag if $|\Delta_{\text{cost}}| > 0.50$ (50% swing).

10.15 Backtest Metrics

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i| \quad (10)$$

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2} \quad (11)$$

$$\text{MAPE} = \frac{1}{n} \sum_{i=1}^n \frac{|y_i - \hat{y}_i|}{|y_i|} \quad (12)$$

$$\text{Dir. Acc.} = \frac{1}{n} \sum_{i=1}^n \mathbf{1}[\text{sign}(\Delta \hat{y}_i) = \text{sign}(\Delta y_i)] \quad (13)$$

Section 11: ML Models—How They Work, Accuracy & How We Got There

11.1 Why These Models?

We chose **statistical / time-series models** rather than deep learning because:

1. **Limited data:** 48 months (4 years) per entity at most—not enough for neural networks
2. **Interpretability:** Scherzinger’s management needs to understand and trust the predictions
3. **Robustness:** Classical models work well with small sample sizes
4. **Transparency:** Every coefficient and weight is explainable

11.2 Model Development Process

1. **Data preparation:** Aggregated invoices into monthly margin series per entity
2. **Quality filtering:** Excluded records with `dq_any_issue = TRUE`
3. **Minimum data requirements:** Each model has minimum month thresholds (12/18/24)
4. **Training:** Models trained on full available history for each entity
5. **Validation:** Walk-forward backtesting with expanding window (never look-ahead bias)

11.3 Accuracy Summary

Model	MAE	RMSE	Dir. Acc.	Best For
Linear Trend	0.024 (2.4pp)	0.028	54.5%	Trend detection
Seasonal	0.032 (3.2pp)	0.037	9.1%	Monthly patterns
Ensemble	0.035 (3.5pp)	0.040	9.1%	Balanced view
EMA	0.065 (6.5pp)	0.068	54.5%	Short-term reaction

Table 12: Model accuracy summary.

What to tell the client:

- “Our best model predicts margin within 2.4 percentage points on average”
- “If we forecast 65% margin, expect the actual to land between 62.6% and 67.4%”
- “The ensemble model balances stability and responsiveness—it’s our recommended default”
- “Directional accuracy is modest (55%) because month-to-month margin is inherently noisy, but the magnitude predictions are very accurate”

11.4 Current Forecast Values

Overall Portfolio:

- Current margin: 59.7%
- 3-month forecast: 60.5% (CI: 57.3%–63.7%)
- 6-month forecast: 60.5% (CI: 57.2%–63.7%)
- 12-month forecast: 60.4% (CI: 57.1%–63.6%)

By Commodity Group (12-month forecast):

Group	Current	12m Forecast	Direction
BKAES	65.0%	63.2%	Slight decline
BKAGG	45.7%	45.5%	Flat (but already low)
BKAIZ	51.6%	56.4%	Improving
MBDIV	41.7%	32.8%	Declining fast
SOPU	58.2%	−45.2%	Collapsing
SOPUZK	80.4%	57.4%	Declining

Table 13: 12-month margin forecasts by commodity group.

Section 12: JSON Data Files—What They Contain

The frontend uses 12 JSON files in `/frontend/src/data/`:

File	Records	Key Contents
<code>dashboard_data.json</code>	1 obj	Annual summaries (4 yrs), monthly revenue (48 pts), commodity breakdown, risk distribution, top 20 customers, quote summary
<code>monthly_detail.json</code>	48 obj	Per-month: revenue, DB1/DB2 margin, invoice count, customer count, avg invoice value
<code>customers_detail.json</code>	25 obj	Top 25 customers: segment, revenue by year, margin by year, top products, win rate, risk score, LTV
<code>products.json</code>	50 obj	Top 50 SKUs: revenue/margin by year, cost breakdown, trend, at_risk flag
<code>forecasting.json</code>	1 obj	Overall + 6 commodity forecasts, model accuracy, seasonal indices, Monte Carlo, backtests
<code>pricing_analysis.json</code>	1 obj	Gap analysis, gap by year, catalog vs quoted, win rate by margin band, rejection codes, t-test
<code>pipeline.json</code>	1 obj	Pipeline stages, conversion funnel, quarterly pipeline, by-commodity pipeline
<code>ml_analytics.json</code>	1 obj	Churn predictions, margin classification, anomaly detection (33), BCG matrix
<code>sales_transactions.json</code>	60 obj	30 recent invoices + 30 recent quotes
<code>inventory_detail.json</code>	1 obj	Inventory metrics, warehouse data
<code>cogs_detail.json</code>	1 obj	Cost of goods sold breakdowns
<code>price_governance.json</code>	1 obj	Price approval workflows

Section 13: Example SKU Scenario 1: Article 300143 (Declining Margin, At Risk)

13.1 The Product

- **Article ID:** 300143
- **Description:** Innenzahnringpumpe (Internal Gear Ring Pump)
- **Commodity Group:** BKAGG
- **Total 4-Year Revenue:** €447,901
- **Total Units Sold:** 757
- **Risk Status:** AT RISK

13.2 Year-by-Year Performance

Year	Revenue	Units	Avg Margin	Trend
2022	€ 148,912	256	57.0%	Baseline
2023	€ 157,706	266	44.5%	–12.5pp drop!
2024	€ 79,831	133	45.5%	Flat but volume halved
2025	€ 61,453	102	42.8%	Still declining

Table 15: Article 300143 year-by-year performance.

13.3 What Happened?

This product shows a classic **margin erosion + volume decline pattern**:

1. **Margin dropped 14.2pp** from 2022 (57.0%) to 2025 (42.8%)
2. **Volume dropped 60%** from 256 units (2022) to 102 units (2025)
3. **Revenue dropped 59%** from € 148K to € 61K

13.4 Cost Breakdown

Component	Share	What It Means
Material	47.8%	Nearly half is raw materials
Direct Mfg (FEK)	27.5%	Significant labor/machine cost
Outsourcing (FV)	3.3%	Low outsourcing dependency
Overhead	21.4%	Fixed cost allocation

13.5 Why It's At Risk

The BKAGG commodity group as a whole is struggling: group-wide margin only 53.8%, growth rate –9.6% (BCG: Dog), Monte Carlo probability below 50%: **61%**. The high material cost share (47.8%) makes it vulnerable to raw material price increases.

13.6 What the Platform Tells Scherzinger to Do

- **Phase 1 insight:** Product is below commodity group benchmark and declining.
- **Phase 2 insight:** Forecast projects continued margin pressure.
- **Phase 3 visibility:** “AT RISK” badge, red color, declining trend arrow.
- **Potential action:** Re-negotiate material pricing, increase customer price, or consider discontinuing.

13.7 Revenue Impact if Nothing Changes

If margin declines at –3.5pp/year and volume drops –50%/year:

- **2026 projected:** ~51 units × € 480/unit = € 24,500 at 39% margin → only € 9,555 DB2
- **vs. 2022:** 256 units × € 582/unit = € 148,912 at 57% → € 84,880 DB2
- **Margin loss over time: € 75,325 per year on this single SKU**

Section 14: Example SKU Scenario 2: Article 201885 (Healthy & Growing)

14.1 The Product

- **Article ID:** 201885
- **Description:** Elektro-Zahnradpumpe (Electric Gear Pump)
- **Commodity Group:** BKAES
- **Total 4-Year Revenue:** €428,847
- **Total Units Sold:** 224
- **Risk Status:** NOT AT RISK—Margin trend: Rising

14.2 Year-by-Year Performance

Year	Revenue	Units	Avg Margin	Trend
2022	€89,561	62	72.1%	Baseline
2023	€150,400	67	77.8%	+5.7pp improvement
2024	€42,208	20	74.1%	Volume dip (margin still good)
2025	€146,678	75	74.5%	Strong recovery

Table 16: Article 201885 year-by-year performance.

14.3 What's Going Right?

1. **Margin improved 2.4pp** from 72.1% to 74.5%, with a peak at 77.8% in 2023
2. **Volume recovered strongly** to 75 units in 2025 after a dip in 2024
3. **Revenue bounced back** to €147K—second-highest year
4. **High-value product:** Avg price/unit €1,914 (vs. €592 for article 300143)

14.4 Cost Breakdown

Component	Share	What It Means
Material	52.3%	Higher but offset by pricing power
Direct Mfg (FEK)	21.0%	Moderate labor cost
Outsourcing (FV)	2.3%	Almost no outsourcing
Overhead	24.4%	Standard allocation

14.5 Why It's Healthy

BKAES is the star performer: 66.8% group margin, +2.7% growth, Monte Carlo below-50% probability only 16.9%. This SKU outperforms its group by ~8pp (74.5% vs 66.8%).

14.6 Revenue Potential

- Current trajectory: ~€150K/year at 74% margin = **€111K annual DB2**
- 20% volume increase: ~€180K/year = **€133K annual DB2**
- **€22K additional annual profit from a single SKU**

14.7 Comparison: SKU 300143 vs 201885

Metric	300143 (At Risk)	201885 (Healthy)
Commodity Group	BKAGG (Dog)	BKAES (Star)
2025 Margin	42.8%	74.5%
Margin Trend	−14.2pp	+2.4pp
Volume Trend	−60% decline	+21% growth
Avg Price/Unit	€ 602	€ 1,956
Risk Status	AT RISK	Safe
Action	Reprice/discontinue	Invest and grow

Table 17: Head-to-head comparison of the two example SKUs.

Section 15: How Much Does This Save Them?

15.1 Direct Savings from Margin Visibility

Before the platform: Scherzinger’s management couldn’t see which customers or products were silently eroding margin. They reviewed aggregate spreadsheets annually.

Margin Gap Recovery: The platform identified a consistent 1.9pp gap between quoted and actual margins. On €6.25M annual revenue, closing even half = **€59,375/year**.

At-Risk Product Identification: 13 products have negative margins. Fixing or discontinuing the worst 5: estimated **€50,000–100,000/year** loss avoidance.

Customer Risk Early Warning: 22 critical-risk customers with €1.14M total LTV. Early intervention could save **€50,000–80,000/year**.

Quote Win Rate Optimization: If strategic repricing improves win rate by 2pp: ~26 additional wins × €7,121 avg deal = **€185,000 additional revenue/year**.

15.2 Total Estimated Annual Impact (Phases 1–3)

Impact Area	Conservative	Optimistic
Margin gap recovery	€ 30,000	€ 60,000
Negative margin product fixes	€ 50,000	€ 100,000
Customer churn prevention	€ 50,000	€ 114,000
Win rate improvement	€ 100,000	€ 185,000
Total Annual Impact	€ 230,000	€ 459,000

Table 18: Estimated annual financial impact.

15.3 Time Savings

Before: Margin analysis done manually in Excel, quarterly, taking 2–3 days per analyst. After: Real-time dashboard refreshes—what took 3 days now takes 30 seconds. Estimated analyst time saved: **~40 hours/quarter = 160 hours/year**.

Section 16: Future Phases—What They Can Expect

16.1 Phase 4: Pricing & FX Analysis (Planned)

1. **Price Optimization Engine:** Analyzes the relationship between price level and win probability. Recommends optimal price points that maximize expected profit (margin $\times$ win probability). Uses the price sensitivity data we already have ($p = 0.006$).
2. **FX Impact Analysis:** Models how currency movements affect landed cost and margin. Provides hedging recommendations when FX exposure exceeds thresholds.
3. **Price Governance Framework:** Automated approval workflows: discounts $>5\%$ require manager approval, $>10\%$ require director. Tracks discount frequency by sales rep, customer, product.
4. **What-If Simulator:**
 - “What happens to margin if material costs rise 10%?”
 - “What if we raise prices on BKAGG products by 5%?”
 - “How much volume can we afford to lose if we increase price by 3%?”

16.2 Phase 5: Advanced ML Analytics (Planned)

1. **ML Churn Prediction:** Trained classifier (Random Forest or XGBoost) predicting 12-month churn probability. Features: order frequency, margin trend, volume trend, days since last order, quote win rate. Expected accuracy: 78–85%.
2. **Customer Lifetime Value (CLV) Modeling:**

$$CLV = \sum_{t=1}^T \text{predicted_revenue}_t \times \text{predicted_margin}_t \times \text{retention_prob}^t$$
3. **Demand Classification:** Segments each SKU into demand patterns: Steady, Seasonal, Lumpy, Intermittent, New. Uses coefficient of variation and order frequency analysis.
4. **Advanced Anomaly Detection:** ML-based outlier detection (Isolation Forest) on margin, revenue, and cost patterns. Automatic alerts when metrics deviate from expected patterns.

16.3 Phase 6+ (Potential Future)

- **ERP Integration:** Direct connection to SAP/ERP for real-time data feeds
- **Email/Alert System:** Automated notifications on risk score changes
- **Sales Rep Dashboard:** Individual views per sales team member
- **Supplier Cost Intelligence:** Track supplier price changes
- **Competitive Intelligence Integration:** Feed competitor pricing into the optimizer

Section 17: Future Phase Example Scenarios

17.1 Scenario A: Phase 4—Price Optimization for Article 300143

Current Situation (Phase 1–3): Article 300143: BKAGG internal gear pump, 2025 margin 42.8% declining, win rate for BKAGG $\sim 35\%$, lost to price (PA+PR) $\sim 7.6\%$.

What Phase 4 Would Do: The Price Optimization Engine would build a win probability curve: $P(\text{win}) = f(\text{margin})$.

```

Article 300143 -- Price Elasticity Model:
Current quoted margin: 48%
Current win rate: 31%

Optimal price point:
If we quote at 44% margin -> expected win rate rises to 38%
Expected profit change: (-4pp margin) x (+22% more wins) = NET
    POSITIVE

Recommendation: Reduce quoted margin to 44% for competitive deals
Expected annual impact: +7 additional wins x EUR 600/unit x 15 units
    avg
                        = EUR 63,000 additional revenue
                        = EUR 27,720 additional DB2

```

What-If Simulator would show:

- “If material costs rise 5%: margin drops to 39.8%. Must raise price by 3% to maintain 42.8%.”
- “If we lose Customer 101690: revenue drops €95K. Need 3 replacement customers.”

17.2 Scenario B: Phase 5—ML Churn Prediction for Customer 101690

Current Situation (Phase 1–3): Customer 101690: €1.54M total revenue (top customer), Enterprise segment, risk tier Medium (score 0.35), margin trend declining (57.5% → 53.6%), win rate 35.5%.

```

Customer 101690 -- Churn Risk Assessment:

Features extracted:
- Order frequency: 28/year avg -> dropped to 24 in 2024 -> back to
  30 in 2025
- Days since last order: 14 days (active)
- Margin trend: -3.9pp over 4 years
- Win rate: 35.5% (below 37.6% average)
- Rejection pattern: 2 price-related losses in 2025
- Average deal size: Increasing (EUR 13.6K -> EUR 15.6K)

ML Prediction:
- 12-month churn probability: 18%
- Classification: WATCHLIST

CLV Estimate:
- Predicted annual revenue: EUR 450K -> 430K -> 410K
- Retention probability: 82% -> 78% -> 74%
- Total CLV: EUR 1,020,000

Recommended Action:
- Schedule quarterly business review
- Offer volume discount on top 3 products
- Investigate the 2 price losses

```

17.3 Scenario C: Phase 5—Demand Classification for Production Planning

Current: Article 201885 orders come in bursts—62, 67, 20, 75 units across four years.

Article 201885 -- Demand Pattern Classification:

Monthly order analysis (48 months):

- Months with orders: 32 out of 48 (67%)
- Coefficient of variation: 0.85 (high variability)
- Order size range: 1-15 units per order
- Seasonal pattern: Stronger in Q1 and Q3

Classification: LUMPY DEMAND

Production Planning Recommendation:

- Base plan: 6 units/month
- Flex capacity: +10 units with 4-week notice
- Annual forecast: 75-85 units (based on 2025 trajectory)

Section 18: API Endpoint Reference

18.1 Phase 1 Endpoints (19)

Endpoint	Method	Purpose
/api/v1/stats	GET	Invoice/quote/customer/product counts
/api/v1/margins/summary	GET	Overall revenue, DB2 margin, counts
/api/v1/margins/by-year	GET	Annual revenue + margin
/api/v1/margins/by-customer	GET	Top N customers by revenue
/api/v1/margins/by-product	GET	Top N products by revenue
/api/v1/margins/by-commodity-group	GET	Revenue + margin per commodity group
/api/v1/margins/gap-analysis	GET	Quoted vs actual margin gap
/api/v1/margins/catalog-vs-quoted	GET	Catalog vs quoted comparison
/api/v1/margins/trend	GET	Monthly/quarterly trends
/api/v1/quotes/summary	GET	Win/loss counts, rates
/api/v1/quotes/win-rate-by-year	GET	Annual win rates
/api/v1/quotes/win-rate-by-deal-size	GET	Win rate by revenue band
/api/v1/quotes/win-rate-by-customer	GET	Win rate per customer
/api/v1/quotes/rejection-codes	GET	Loss reasons + revenue impact
/api/v1/quotes/price-sensitivity	GET	T-test: pricing vs win probability
/api/v1/data-quality/summary	GET	Quality percentages
/api/v1/data-quality/issues	GET	Individual DQ records
/api/v1/dashboard	GET	Composite dashboard data

18.2 Phase 2 Endpoints (9)

Endpoint	Method	Purpose
/api/v1/forecasts/{type}/{id}	GET	Margin forecast for entity
/api/v1/forecasts/{type}/{id}/compare	GET	All models side-by-side
/api/v1/forecasts/accuracy	GET	Backtest accuracy by model

/api/v1/risk/scores	GET	Risk scores (filter by tier)
/api/v1/risk/scores/{customer_id}	GET	Risk detail + components
/api/v1/costs/trends	GET	Product cost trends by quarter
/api/v1/costs/risers	GET	Top cost-rising products
/api/v1/benchmarks	GET	All commodity benchmarks
/api/v1/benchmarks/{commodity_group}	GET	Benchmark for specific group
/api/v1/simulations/{type}/{id}	GET	Monte Carlo results

Section 19: Key Numbers to Remember for the Demo

19.1 The Big Numbers

- **€24.6M** total revenue over 4 years
- **5,565** invoices analyzed
- **4,605** quotes analyzed
- **967** unique customers
- **1,223** unique products
- **64.78%** average DB2 margin
- **37.6%** overall win rate (improving: 36.5% → 40.4%)

19.2 The Problem Numbers (What They Need to Fix)

- **€24.9M** in lost quotes (2025 alone)
- **€971K** lost to price competition (PA + PR codes)
- **€370K** lost because they didn't follow up (KD code!)
- **47.6%** of scored customers at High or Critical risk
- **1.9pp** consistent margin erosion from quote to invoice
- **13** products with negative margins (losing money on every sale)
- **BKAGG** (−9.6% growth) and **SOPU** (−96.2% growth) declining
- **MBKUEHL** at only 9.8% margin—nearly unprofitable

19.3 The Opportunity Numbers (What They Can Gain)

- **€230K–459K** estimated annual impact from insights
- **+2.4pp** margin improvement potential from closing the quote–invoice gap
- **26+ additional won deals** if pricing is optimized (2pp win rate improvement)
- **160 hours/year** analyst time saved
- **BKAIZ** growing +70.9%—invest here
- **BKAES** is the reliable star—66.8% margin, 68% of revenue

19.4 The Forecast Numbers

- Overall 12-month forecast: **60.4%** (stable, slight decline from 59.7%)
- Best performing group: **SOPUZK** at 80.4% current
- Best forecast trend: **BKAIZ** improving from 51.6% to 56.4%
- Biggest concern: **SOPU** forecasted to go negative (−45.2% at 12 months)
- Monte Carlo: **25.6% probability** overall margin drops below 50%

19.5 The Model Numbers

- Best model: **Linear Trend**—MAE 0.024 (2.4pp accuracy)
- Risk model accuracy: **78%** (churn prediction)

- Monte Carlo: **10,000 simulations** per entity
- Risk scoring: **5 components**, all weights add to 100%

This document covers all completed Phase 1–3 work, including every data source, transformation, formula, ML model, accuracy metric, JSON file, graph, and API endpoint. The two SKU examples (300143 declining vs. 201885 healthy) demonstrate real patterns from the actual data. Future phase scenarios show the tangible value the client can expect from continued investment.